

Who's Building Agentic Apps — and How

A plain-language map of the companies and tools building software that gives users freedom through AI agents, ranked by scale, with three standout projects examined up close. Written for people who know AI but don't write code.

Compiled 2026-05-26 · Figures gathered from company sites, press, and live GitHub counts · Numbers move fast and several are "reported / in talks" — see footnotes & sources.

★ Start here — the 60-second version

"Agentic apps" are tools where an AI *does multi-step work on your behalf* — not just answering, but acting: writing & running code, browsing the web, resolving support tickets, organizing files. The market split into four buckets:

- **Coding agents** — the biggest commercial winners so far (Cursor, Devin, Replit, Claude Code).
- **Build-your-own frameworks** — open-source toolkits devs use to assemble agents (LangChain, CrewAI, n8n, Dify).
- **Computer/browser-use agents** — AI that drives a real screen (ChatGPT Agent, Manus, Browser Use).
- **Vertical enterprise agents** — agents sold to a specific industry, where the real money is (Sierra=support, Harvey=legal, Glean=search).

If you read nothing else, jump to the [three standouts](#): **Cursor** (fastest-growing software company ever), **Manus** (the surprise viral "general agent"), and **Sierra** (agents that only get paid when they solve your problem).

Open-source badge: OPEN code is public & free to run · CLOSED proprietary SaaS. **Footnote** ¹ = number is reported/in-talks, not officially confirmed.

What's inside

1. The big players, ranked by scale — commercial agentic-app companies by valuation / revenue / users
2. The open-source toolkits, ranked by traction — what builders actually use, by GitHub stars
3. Three standouts up close — Cursor, Manus, Sierra
4. Head-to-head comparison — what each one teaches us
5. What it means for Atram

1 · The big players, ranked by scale

These are the funded companies building agentic products, ordered roughly by company value (with revenue and user signals alongside). Valuations for private companies are point-in-time and several recent ones are still "reported." Treat the order as a tier guide, not a precise league table.

#	Company / Product	What its agent does	Scale (value · revenue · users)	Pricing	Code	Link
1	Cursor CODING	AI code editor that plans & edits across your whole project, runs background agents.	~\$29.3B ¹ valuation · ~\$2B ARR · fastest to \$2B ever	Free / \$20 / \$60 / \$200 mo	CLOSED	cursor.com
2	Perplexity SEARCH / BROWSER	AI answer engine + "Comet" agentic browser that does tasks while you read.	~\$20B valuation · ~\$200M ARR · ~45M monthly users	Free / Pro \$20 / Max \$200	CLOSED	perplexity.ai
3	Sierra ENTERPRISE · CX	Customer-service agents that resolve tickets end-to-end. (Bret Taylor's company.)	~\$15B valuation · ~\$150M ARR · 40%+ of Fortune 50	Pay-per-resolution (outcome-based)	CLOSED	sierra.ai
4	Harvey ENTERPRISE · LEGAL	Legal AI agents for contracts, research, diligence at law firms.	~\$11B valuation · ~\$190M ARR · 100k+ lawyers	Enterprise / per-seat	CLOSED	harvey.ai
5	Notion PRODUCTIVITY	Custom AI agents that run tasks inside your workspace docs & databases.	~\$11B valuation · ~\$600M ARR	\$20/user + agent credits	CLOSED	notion.so
6	Cognition (Devin + Windsurf) CODING	"Devin," an autonomous software-engineer agent; also	~\$10.2B ¹ (talks of ~\$25B ¹) · Devin ~\$73M ARR	\$20 / \$500 / mo + enterprise	CLOSED	devin.ai

owns the Windsurf IDE.

7	Replit CODING · NO-CODE	Cloud workspace where an agent builds & ships whole apps from a prompt.	~\$9B valuation · ~\$150M ARR · 50M+ users	Free / \$20 / \$100 mo	CLOSED	replit.com
8	Glean ENTERPRISE · SEARCH	Enterprise search + agents that act across all your work apps.	~\$7.2B valuation · ~\$200M ARR	Enterprise per-seat + usage	CLOSED	glean.com
9	Decagon ENTERPRISE · CX	"Concierge" customer-support agents.	~\$4.5B valuation · ~\$35M ARR · 100+ logos	Enterprise (outcome/usage)	CLOSED	decagon.ai
10	n8n AUTOMATION	Visual workflow automation with native AI agents (400+ integrations).	~\$2.5B valuation · 190k★ on GitHub	Self-host free / Cloud €24+ mo	OPEN*	n8n.io
11	Manus GENERAL AGENT	A "do the whole task" autonomous agent on its own virtual computer.	Reported acquired by Meta at ~\$2B+ ¹ (Dec 2025) · ~\$100M ARR in 8 mo	Credit-based subscription	CLOSED	manus.im
12	Factory CODING	"Droids" — agents that handle the full software lifecycle.	~\$1.5B valuation · used at Nvidia, Adobe	Free / \$20 mo / enterprise	CLOSED	factory.ai
13	Genspark GENERAL AGENT	"Super agent" that builds deliverables and even makes phone calls.	~\$1.25B valuation · \$36M ARR in 45 days	Freemium + credits	CLOSED	genspark.ai
14	LangChain FRAMEWORK CO.	The most-used toolkit for building agents; sells hosting & monitoring.	~\$1.25B valuation · ~90M monthly downloads	OSS free / cloud usage-based	OPEN	langchain.com
15	Augment Code CODING	Enterprise coding agent with a huge codebase "context engine."	~\$1B valuation · ~\$20M ARR ¹	Credit-based + enterprise	CLOSED	augmentcode.com
16	Gumloop AUTOMATION	No-code builder that turns any employee into an agent-builder.	\$50M Series B (Benchmark) · Shopify, Ramp	Free / \$37 / \$244 mo	CLOSED	gumloop.com

Also enormous, but baked into trillion-dollar companies (so not separately ranked): **OpenAI ChatGPT Agent** (absorbed "Operator"), **Anthropic Claude** (computer use + Claude Code), **Google Gemini / Project Mariner**, and **Microsoft Copilot**. These are arguably the most-used agents of all — they just aren't standalone startups. Smaller but notable: **Lindy & Relay.app** (no-code agents), **Cora & Sparkle** (Every's consumer apps), **Hebbia** (finance/legal).

2 · The open-source toolkits, ranked by traction

These are mostly free, open-source building blocks that developers use to assemble their own agents. "Stars" are GitHub's bookmark/popularity signal (fetched May 2026) — a good proxy for mindshare, though older projects can carry legacy hype. Most are free to use; the companies behind them sell hosting, monitoring, or enterprise tiers.

#	Project	What it is	GitHub ★	How it makes money	License	Link
1	n8n	Visual AI-workflow automation	~190k	Cloud tiers / enterprise	Fair-code	repo
2	AutoGPT	The original "autonomous GPT"; now a low-code platform	~185k	Cloud platform (beta)	MIT*	repo
3	Diffy	Build & run agentic workflows (visual)	~143k	Cloud \$59–159 mo	Apache*	repo
4	LangChain	General agent-building framework	~138k	LangSmith / LangGraph cloud	MIT	repo
5	Claude Code ANTHROPIC	Terminal coding agent	~127k	Via Claude plans / API	Proprietary	repo
6	Gemini CLI GOOGLE	Open terminal AI agent	~105k	Free tier / API	Apache-2.0	repo
7	Browser Use	Let agents control a web browser	~96k	Cloud API (freemium)	MIT	repo
8	OpenAI Codex OPENAI	Terminal + cloud coding agent	~86k	Via ChatGPT / API	Apache-2.0	repo
9	OpenHands	Open autonomous dev-agent platform	~75k	Cloud + enterprise	MIT	repo
10	MetaGPT	Multi-agent "AI software company"	~68k	Atoms / MGX product	MIT	repo
11	Cline	Autonomous coding agent (VS Code)	~62k	Bring-your-own model keys	Apache-2.0	repo
12	AutoGen MICROSOFT	Multi-agent framework (<i>now maintenance-only</i>)	~58k	Free (Azure upsell)	MIT	repo
13	CrewAI	Orchestrate role-playing agent "crews"	~52k	Enterprise tier	MIT	repo
14	LlamaIndex	Document/RAG agent framework	~50k	LlamaCloud credits	MIT	repo
15	Goose BLOCK	Local, extensible AI agent (MCP-first)	~46k	Free	Apache-2.0	repo
16	Aider	Terminal AI pair-programmer	~45k	Free (your API costs)	Apache-2.0	repo

Other strong toolkits: **Agno** (~40k), **LangGraph** & **Continue** (~33k), **Composio** (~29k, tool-connections), **Smolagents** (HuggingFace, ~28k), **OpenAI Agents SDK** (~27k), **Vercel AI SDK** & **Mastra** (~24k, TypeScript), **Letta** (~23k, agent memory), **Skyvern** (~22k, browser RPA), **Google ADK** (~20k), **E2B** (~12k, secure sandboxes). * "Fair-code"/modified licenses aren't fully open-source in the strict sense.

3 · Three standouts, up close

Chosen for variety and significance: the breakout commercial success, the biggest surprise, and the most forward-looking business model.

Cursor — the breakout success MOST SUCCESSFUL

Anysphere · AI code editor · closed-source

Cursor is an AI-first code editor where an agent reads your whole project, proposes multi-file changes, and runs "background agents" that work while you do something else. It became the **fastest software company in history to reach \$2B in annual revenue**, reportedly going from ~\$100M to ~\$2B ARR in about a year.

~\$29.3B valuation¹ ~\$2B ARR \$20–200/mo

Why it matters: Cursor is the proof that agentic apps are a real, enormous market — not hype. It also shows the winning consumer playbook: take an everyday workflow (writing code), make the agent feel like a natural extension of it, and charge a simple subscription.

Manus — the surprise MOST SURPRISING

Butterfly Effect / Monica (China) · general autonomous agent · closed-source

Manus went viral in March 2025 as a "general agent" that takes a goal, spins up its **own virtual computer**, and completes the entire multi-step task — research, building files, booking things — then hands you the finished result. It reportedly hit **~\$100M revenue in about 8 months** (among the fastest ever) and was **reportedly acquired by Meta in Dec 2025 for \$2B+¹**.

~\$100M ARR in 8mo reported **Meta** acq. \$2B+¹ 2M+ week-1 waitlist

Why it matters: Manus is the clearest bet on the "do the whole job for me" future — and a reminder that a breakout agentic app can now come from anywhere and scale in months. (The Meta deal price is press-reported, not officially confirmed.)

Sierra — the forward-looking model MOST FORWARD-LOOKING

Bret Taylor (ex-Salesforce co-CEO, OpenAI chair) · enterprise customer-experience agents · closed-source

Sierra builds customer-service agents that resolve issues end-to-end for big companies. Its standout idea is **outcome-based pricing**: instead of charging per seat or per message, Sierra largely charges **per problem actually resolved**. It reached **~\$150M ARR** and a **~\$15B valuation** in under two years, serving 40%+ of the Fortune 50.

~\$15B valuation ~\$150M ARR pay-per-resolution

Why it matters: "Pay only when the agent succeeds" may be the pricing model the whole industry moves toward — it aligns cost with value and is only possible because the agent does real work, not just assist. Worth watching for how we'd ever price an Atram agent.

Honorable mentions: **Devin** (the autonomous-engineer bet), **Browser Use** (a tiny open-source project at ~96k stars powering computer-use everywhere), and **MCP** — Anthropic's open "plug" standard that lets any agent connect to any tool (now an industry default).

4 · Head-to-head — what each teaches us

	Cursor	Manus	Sierra
Who it's for	Developers (millions)	Anyone, any task (consumer + pro)	Big enterprises (their customers)
What the agent does	Writes & edits code in your project	Completes whole tasks on a virtual PC	Resolves customer-support cases
How you pay	Flat monthly subscription	Subscription + usage credits	Per resolved outcome
Scale	~\$2B ARR · ~\$29.3B ¹	~\$100M ARR · reported Meta acq. ¹	~\$150M ARR · ~\$15B
Its big idea	Make the agent native to a daily workflow	One agent for <i>everything</i>	Charge for results, not seats

Main risk	Frontier labs (OpenAI/Anthropic) competing directly	"Generalist" agents may lose to focused ones; reliability	Hard to guarantee outcomes; long enterprise sales
-----------	---	---	---

The through-line: **the winners pick a clear job, let the AI do it end-to-end, and price around the value delivered**. Cursor proves the market is huge, Manus shows breakout speed is possible from anywhere, and Sierra shows where pricing is heading. Note all three are **closed-source products** — the open-source world (Section 2) is mostly the *picks-and-shovels* layer that everyone, including these companies, builds on.

5 · What it means for Atram

- 1. The market is proven and fast.** Agentic apps went from curiosity to multi-billion-dollar businesses in ~18 months. Speed-to-traction is unusually high — the bar for "interesting" is now "does the whole job," not "assists with it."
- 2. Pick a job, own it end-to-end.** Every winner is built around one clear outcome (write code / resolve a ticket / do the task). For our own products, the lesson is to define the outcome sharply and let the agent own the full loop — the Every "agent-native" idea, validated at scale.
- 3. Watch outcome-based pricing.** Sierra's "pay when it works" model is the most forward-looking commercial idea here and likely where client expectations head next.
- 4. Build on the open layer.** We don't need to reinvent the plumbing — LangChain, the MCP standard, Browser Use, and sandbox runtimes (E2B) are mature, free building blocks. Buy/borrow the picks-and-shovels; spend our effort on the specific outcome we deliver.

Sources, dates & honesty notes

Compiled 2026-05-26 from company sites, reputable press (TechCrunch, CNBC, Bloomberg, VentureBeat), funding announcements, and GitHub star counts fetched live the same week. **Valuations and revenue for private companies are point-in-time and move quickly.** Footnote 1 marks figures that were **reported or "in talks" rather than officially confirmed** — specifically: Cursor's ~\$29.3B (and a rumored ~\$50B round), Cognition's ~\$10.2B (and rumored ~\$25B), Manus's reported Meta acquisition at \$2B+, and Augment's ~\$20M ARR. GitHub stars are popularity signals, not usage. A companion **Technical Edition** of this report covers architectures, protocols (MCP/A2A), benchmarks (SWE-bench, OSWorld), and current research. Lives in `.github/docs/claude-code/`. Questions → Engineering channel.