

How Every Builds Agentic Apps

What the Every team (the makers of the **compound-engineering** plugin many of us run in Claude Code) believes about building software that hands real freedom to the user through LLMs — plus every app they ship, where to get it, and what's worth your time.

Researched & compiled 2026-05-25 · Sources verified against every.to, the EveryInc GitHub org, and live search · Caveats at the end.

★ Start here if you're a PM or client-facing (not a dev)

You don't need to read this whole thing. Here's the 60-second version and the three places worth your time:

The big idea: Every (the AI studio led by Dan Shipper, with Kieran Klaassen behind the compound-engineering plugin) argues the next generation of apps shouldn't hard-code every button and flow. Instead you give an LLM a few simple tools and let it figure out *how* to get the user's outcome — so users change what the app does just by **asking in plain language**. They call this "**agent-native**."

- **Read one essay:** "Agent-native Architectures: How to Build Apps After the End of Code" — see [Section 1](#). It's written for thinkers, not coders.
- **Try one app:** their products in [Section 2](#) are the idea made real — **Cora** (email that triages itself), **Sparkle** (a Mac that cleans itself when you ask), **Monologue** (dictation that fixes your tone).
- **Listen on a commute:** the talks & podcast in [Section 4](#) explain all of this conversationally — no code required.

Badges used in this guide: **FOR EVERYONE** approachable, no code needed **MORE TECHNICAL** aimed at engineers **TRY / INSTALL** something you can actually use

What's inside

1. The Big Ideas — why apps should give users freedom through LLMs (the essays)
2. The Apps — everything Every ships, with links, what it does, and who it's for
3. The Compound-Engineering Plugin — the one in our Claude Code, and how to install it
4. Watch, Listen & Follow — podcast episodes, talks and accounts (great for non-devs)
5. Takeaways for Atram — why this matters for how we build

1 · The Big Ideas — software that gives users freedom through LLMs

The core thesis across Every's writing: stop hard-coding every feature. Define a few atomic tools, let an agent decide how to combine them, and let the user steer it in plain language. These are the essays that make the case, most useful first.

Agent-native Architectures: How to Build Apps After the End of Code **FOR EVERYONE**

Dan Shipper · every.to

The flagship piece. The core of an app shouldn't be code but an agent — "something squishy and alive." Developers describe the *what* (the outcome) and let the agent work out the *how*. The payoff for users: they can change how the app behaves **just by changing words in a language we already speak**.

Why it matters for you: this is the single best explainer of the whole philosophy, written for a general audience. If a client asks "why is your AI product different from a normal app?" — this is your answer.

Read: every.to/chain-of-thought/agent-native-architectures... · hub: every.to/guides/agent-native

How to Build Agent-native: Lessons From Four Apps **FOR EVERYONE**

Katie Parrott · every.to

The practical companion. Three design principles in plain terms: **parity** (anything the user can do, the agent can do too), **granularity** (give the AI small, simple tools), and **composability** (small tools combine in ways nobody explicitly designed). Key honesty: hard guarantees like "confirm before deleting" must live in real code — "prompts are still suggestions."

Why it matters for you: it walks through four real products (Cora, Sparkle, Monologue, a book app), so you see the abstract idea become shippable features. Great for grounding a product conversation.

Read: every.to/source-code/how-to-build-agent-native-lessons-from-four-apps

The Knowledge Economy Is Over. Welcome to the Allocation Economy. **FOR EVERYONE**

Dan Shipper · every.to

The worldview behind it all: as AI does more knowledge work, every worker shifts from *maker* to *manager/allocator* — you direct AI toward outcomes instead of executing each task. "Every maker becomes a manager."

Why it matters for you: this is basically a description of the modern PM/operator role. Useful framing for how our own teams (and clients' teams) will work alongside agents.

Read: every.to/chain-of-thought/the-knowledge-economy-is-over...

The Folder Is the Agent **MORE TECHNICAL**

Kieran Klaassen · every.to

Argues you don't need heavy multi-agent frameworks. A plain project folder — an instructions file plus accumulated context — *is* an agent: point a general model at it and it becomes a domain specialist. Freedom comes from **context**, not more code.

Why it matters for you: lighter read for the curious. Explains why our own Claude Code setups get smarter the more context we feed them.

Read: every.to/source-code/the-folder-is-the-agent

Compound Engineering: How Every Codes With Agents **MORE TECHNICAL**

Dan Shipper & Kieran Klaassen · every.to

How a tiny team ships like a big one: nobody writes code by hand; engineers orchestrate parallel agents in a loop — **plan** → **work** → **review** → **compound** — spending ~80% on planning/review and ~20% on coding, so each feature makes the next one *easier*.

Why it matters for you: this is the method behind the plugin in our Claude Code. Helps you understand why "planning and review" is where the real work happens now.

Read: every.to/chain-of-thought/compound-engineering-how-every-codes-with-agents · hub: every.to/guides/compound-engineering

More for the curious **MORE TECHNICAL**

Various · every.to

Deeper cuts worth bookmarking:

- **Compound Engineering: The Definitive Guide** — every.to/source-code/compound-engineering-the-definitive-guide
- **My AI Had Already Fixed the Code Before I Saw It** (the origin essay) — every.to/source-code/my-ai-had-already-fixed-the-code-before-i-saw-it
- **Compound Engineering Camp: Every Step, From Scratch** (1-hour build walk-through) — every.to/source-code/compound-engineering-camp-every-step-from-scratch
- **The Magic Minimum for AI Agents** — every.to/chain-of-thought/the-magic-minimum-for-ai-agents

2 · The Apps — everything Every ships

Every runs a software studio; its full lineup lives at every.to/studio, and most apps are bundled in **The Every Bundle** (~\$30/mo). Most consumer apps are closed-source SaaS, but several have open-source companion tools (CLIs, MCP servers, SDKs) under github.com/EveryInc. The **OPEN** / **CLOSED** pill shows whether there's code you can read.

App	What it does & why it's "agentic"	Get it / links	Good for...
Cora CLOSED	AI email assistant — screens your inbox, batches & triages mail, drafts replies, sends a daily brief. Agentic but deliberately cautious: never sends or deletes without your OK.	Site: cora.computer Sign up: cora.computer/users/sign_up (7-day trial)	Anyone drowning in email. The clearest "agent does the work for you" demo to show a client.
Sparkle OPEN*	AI Mac file organizer — finds junk & duplicates and tidies your files from a plain-language instruction ("say it your way") instead of manual rules.	Site: makeitsparkle.co Docs: help.every.to/.../sparkle Open MCP server: github.com/EveryInc/sparkle-mcp-server	Mac users with messy Downloads folders. *App is closed; the MCP connector is open source.
Monologue OPEN*	AI voice dictation — turns speech into polished text in your tone, drops filler words, 100+ languages, writing modes (email/notes/code).	Site: monologue.to App Store (Mac/iPhone/iPad/Watch): apps.apple.com/.../id6755956193 Open toolkit: github.com/EveryInc/monologue-toolkit	Anyone who'd rather talk than type. Free tier available. *Companion CLI is open.
Spiral CLOSED	AI writing partner "with taste" — turns one idea into social posts / marketing copy / long-form, asks clarifying questions and matches your style.	Site: writewithspiral.com Web app: app.writewithspiral.com Docs: help.every.to/.../spiral	Marketing & client-comms folks who repurpose content. CLI exists: <code>npm i -g @every-env/spiral-cli</code>
Proof OPEN	Collaborative doc editor where humans and AI agents co-author — live agent presence, suggestions, and character-level "who wrote this" provenance. The purest example of the philosophy.	Site: proofeditor.ai Start: proofeditor.ai/get-started Open SDK: github.com/EveryInc/prof-sdk	Free, no login. Devs who want a reference implementation they can actually read.
Plus One CLOSED	One-click personal AI agent inside Slack — connected to Every's apps + your tools, with memory, personality, and recurring scheduled work.	Waitlist: every.to/plus-one (paid subscribers prioritized)	Early-adopters curious about a persistent personal agent. Not GA yet.

Also on the studio page under **Experiments** (Early-Adopter only, no public repos): **Kairos**, **TLDR**, **Extendable Media**, **MindTune**, **Machcast**. Note: **Lex** (lex.page) was incubated at Every but spun out as its own company — it's no longer an Every product. We found no public Every product called "Diane."

3 · The Compound-Engineering Plugin **TRY / INSTALL**

This is the plugin many of us already run in Claude Code — the source of the `/plan`, `/work`, `/review`, `/brainstorm` commands and all those review agents. It's free, open source (MIT), by Kieran Klaassen at Every.

Repository: github.com/EveryInc/compound-engineering-plugin

Install in Claude Code (paste these two lines into Claude Code):

```
/plugin marketplace add EveryInc/compound-engineering-plugin
/plugin install compound-engineering
```

Also works in Cursor, GitHub Copilot CLI, Factory Droid, Qwen Code and Codex (the repo README has the exact command for each).

What you get: a guided workflow — **brainstorm** → **plan** → **work** → **review** → **compound** — plus dozens of specialist "review agents" (security, performance, architecture, framework-specific). The whole point: it pushes ~80% of effort into planning and review, where humans add the most value, and it *remembers* lessons so the next task is easier.

Heads-up for whoever maintains our setup: the version installed on this machine (v2.34.0) is older than the current GitHub `main` (more agents/skills, renamed `ce-` commands). Run `/plugin marketplace update` to refresh.

4 · Watch, Listen & Follow FOR EVERYONE

The easiest way in if you don't want to read essays. All of this explains the philosophy conversationally.

AI & I podcast (Dan Shipper)

Hub: every.to/podcast

- **How Two Engineers Ship Like a Team of 15 With AI Agents** — every.to/podcast/how-two-engineers-ship-like-a-team-of-15-with-ai-agents
- **The AI Sandwich: Where Humans Excel in an AI World** — every.to/podcast/transcript-the-ai-sandwich...
- **How We Built Cora (behind the scenes)** — every.to/podcast/transcript-how-we-built-our-ai-email-assistant...

► **Talks & interviews (YouTube)**

- Kieran — **How to Make Claude Code Better Every Time You Use It** (50-min tutorial) — youtube.com/watch?v=g6z_4TMDiaE
- Kieran — **Compound Engineering: Manage Teams of AI Agents** — youtube.com/watch?v=srh0zy1MQcl
- Dan — **How to be "Agent Native" in 2026** (The Neuron) — youtube.com/watch?v=rHzN8ul9cn0
- Dan — **The AI-native startup: 5 products, 100% AI-written code** (Lenny's) — lennysnewsletter.com/p/inside-every-dan-shipper

Accounts to follow

- **Every** — site every.to · X [@every](#) · YouTube [@everyinc](#)
- **Dan Shipper** (CEO) — X [@danshipper](#) · [LinkedIn](#)
- **Kieran Klaassen** (compound engineering, GM of Cora) — X [@kieranklaassen](#) · [LinkedIn](#)

5 · Why this matters for Atram

1. The product lesson. When we build AI features, the agent-native pattern (a few atomic tools + an LLM deciding how to combine them + the user steering in plain language) often beats hard-coding every flow. Users get freedom; we ship faster. **Cora** and **Sparkle** are the demos to point clients at.

2. The safety caveat. Every is explicit: *prompts are suggestions, not guarantees*. Anything that must always happen — confirmations, permissions, "never send without approval" — belongs in real code, not in a prompt. Worth repeating in any client conversation about reliability.

3. The way-of-working lesson. The compound-engineering loop (plan → work → review → compound, 80% planning/review) is already in our Claude Code. The mindset shift for everyone — devs and PMs — is that **planning and reviewing are now the high-value work**, and the system should get smarter every time we use it.

Sources & honesty notes

All links were verified against **every.to**, the **EveryInc** GitHub org, and live search on 2026-05-25. The install commands and repo are confirmed (the plugin is already installed on the compiling machine). A few items carry small caveats: individual X/Twitter posts and Every's exact LinkedIn company URL came from search-result listings (X profiles were not opened directly); product pricing can change; and "Experiment" apps have no public repos. Treat product pages as the source of truth for current pricing.

Compiled for Atram PMs, client-facing colleagues, and engineers who use Claude Code. Lives in `.github/docs/claude-code/` alongside the Claude Code reference. Questions → ping the Engineering channel.